

Assignment 3 Part B

Group -11

Abhanshu Gupta (agupta89@wisc.edu)

Parikshit Sharma (psharma43@wisc.edu)

Vishnu Lokhande (lokhande@wisc.edu)

Application 1

Pagerank application was developed using graphX api's. Two key Api's from graphX which were used were (a) outerJoinVertices and (b) aggregateMessages.

This application has been compared with Assignment1-PartC-Question3 where just spark api's were used to run the pagerank algorithm.

Inorder to make a uniform comparison, the configurations parameters were set the same in both the applications. Specifically, these were the configuration parameters.

```
("spark.locality.wait", "0")
("spark.driver.memory", "1g")
("spark.executor.memory", "18g")
("spark.executor.instances", "5")
("spark.executor.cores", "4")
("spark.task.cpus", "1")
("spark.serializer", "org.apache.spark.serializer.KryoSerializer")
("spark.logLineage", "false")
```

The output of the application is saved on hdfs at /pagerank_results. The application

	<i>Pagerank (spark api's)</i>	<i>Pagerank (graphX api's)</i>	completion time is measured by the time taken by the application from the beginning of reading time to saving the
<i>Time Taken</i>	29 min	5.4 min	
<i>Network Read</i>	13242 MB	18344 MB	
<i>Network Write</i>	13549 MB	19026 MB	
<i>Storage Read</i>	1434 MB / 7 MBps	1368 MB / 9 MBps	
<i>Storage Write</i>	24249 MB / 2 MBps	4807 MB / 2 MBps	
<i>Number of Stages</i>	24	68	
<i>Number of Tasks</i>	480	544	

results on hdfs. Here is summary.

Comparison

Additional Benefits of graphX:

GraphX provides api's to work directly with the edge and vertex attributes of a graph. This makes developing a graph based application lot easier. One example of this is the "aggregateMessages" api, which passes messages from previous vertices to next vertices over the edges.

GraphX provided "TripletFields.Src" functionality, which helps populate only the necessary fields of an RDD/graph thereby further improving the efficiency.

GraphX when used for graph applications provides better visualizations in the form of triplet views, vertex views and edge views. This is not the case for spark as there is only RDD based view.

graphX vs. Spark

It is observed that pagerank using graphX takes lot less time compared to that of Spark. This is possibly because the ranks graph is being cached in every iteration in the graphX application unlike the spark application. Ranks are used in every iteration, hence, caching them seems to improve the performance a lot. We did not cache ranks in the spark application because we did not get much improvements then as we were dealing with small size datasets. Another possible reason might be that the outerJoinVertices api used in graphX might be more optimized than that of graphX. The number of tasks in spark are less than that of graphX. The number of tasks of an application usually dependent on the level of parallelism which is related to the number of partitions in the application. We used 20 paritions in spark app (because we used 20 paritions before) as compared to 8 paritions in graphX app (default for the datasize).

Application 2

Input: For this application input is the output file "Question2_words.txt" of PartB from Assignment2. We have written a python script "language_cleanup.py" to cleanup web addresses, hashtags and non-alphanumeric words and stored it in the new file "graph_input.txt". This file will serve as input for all sub-applications that are developed under Application2.

Question1

Implementation: PartBApplication2Question1.scala

Logic:

1. Read the "graph_input.txt" file and create a VertexRDD which is a RDD(VertexId,Array[String]) using 'map' function.
2. Create EdgeRDD by doing cartesian product of vertexRDD with itself and then filter out the entries where there is no common word between two vertices
3. Create a graph using VertexRDD and EdgeRDD.
4. Get the triplet view of the graph and use filter on the condition that source word list should be larger than the destination word list.
5. Get the count of the filtered RDD as answer

Output: For my case of 262 vertex the total number of edges that satisfy the criteria: 33648

Question2

Implementation: PartBApplication2Question2.scala

Logic:

1. Read the "graph_input.txt" file and create a VertexRDD which is a RDD(VertexId,Array[String]) using 'map' function.
2. Create EdgeRDD by doing cartesian product of vertexRDD with itself and then filter out the entries where there is no common word between two vertices.
3. Create a graph using VertexRDD and EdgeRDD.
4. Get the aggregated message of counter from each neighbor, thereby creating a neighbor RDD which stores the number of neighbors of each vertex.
5. Find the maximum degree in the created RDD.
6. Filter the RDD to only have entries corresponding to maximum degree.
7. If RDD has only one entry than the entry is the answer.
8. Else find the entry which has maximum words and that entry will be the answer.

Output: For my case of 262 vertex the vertex most popular corresponds to time interval: 53

Question3

Implementation: PartBApplication2Question3.scala

Logic:

1. Read the "graph_input.txt" file and create a VertexRDD which is a RDD(VertexId,Array[String]) using 'map' function.
2. Create EdgeRDD by doing cartesian product of vertexRDD with itself and then filter out the entries where there is no common word between two vertices.
3. Create a graph using VertexRDD and EdgeRDD.

4. Get the aggregated message of word size from each neighbor and reducing it by adding, thereby creating a neighbor RDD which stores the total number of words in neighbors of each vertex.
5. Use map function to take average value on each entry of RDD.
6. Print the RDD values.

Output: The output is list of tuples of the form (VertexId,Average Words in neighbor), printed one vertex per line.

Question4

Implementation: PartBApplication2Question4.scala

Logic:

1. Read the "graph_input.txt" file and create a VertexRDD which is a RDD(VertexId,Array[String]) using 'map' function.
2. Flatten the vertexRDD to get list of words and then reduce them on key to get words and their count.
3. Find the maximum count in the created RDD.
4. Print any entry randomly from the RDD that has its count equal to maximum count.

Output: For my case of 262 vertex the word most popular comes out to be: rt

Question5

Implementation: PartBApplication2Question5.scala

Logic:

1. Read the "graph_input.txt" file and create a VertexRDD which is a RDD(VertexId,Array[String]) using 'map' function.
2. Create EdgeRDD by doing cartesian product of vertexRDD with itself and then filter out the entries where there is no common word between two vertices.
3. Create a graph using VertexRDD and EdgeRDD.
4. Call the connectedComponents() algorithm on this graph to get connected components.
5. Reduce the RDD based on cluster's min vertex and use a counter to keep track of nodes.
6. Print the RDD value that has the maximum count as the biggest cluster.

Output: For my case of 262 vertex the biggest size of subgraph comes to be: 262

Question6

Implementation: PartBApplication2Question6.scala

Logic:

1. Read the "graph_input.txt" file and create a VertexRDD which is a RDD(VertexId,Array[String]) using 'map' function.

2. Flatten the vertexRDD to get list of words and then reduce them on key to get words and their count.
3. Find the maximum count in the created RDD.
4. Locate a word in the RDD that has maximum count.
5. Filter and print all the vertex in vertexRDD that have the popular word in its word list.

Output: For my case of 262 vertex the all the 262 vertex has the most popular word rt