

Assignment 3 Part A

Group -11

Abhanshu Gupta (agupta89@wisc.edu)

Parikshit Sharma (psharma43@wisc.edu)

Vishnu Lokhande (lokhande@wisc.edu)


Machine Learning in Tensor Flow versus Spark

Apache Spark is a general-use data processing engine with libraries that make it easy to do machine learning, batch processing, streaming, and graph computation. Tensorflow is a library used primarily for more scalable and easier machine learning.

- Tensor Flow code can be executed on a variety of devices from clusters to mobile devices, this makes it very extensible.
- TF provides more flexibility in defining on which device to place an operation. It can make use of device heterogeneity e.g. GPUs/TPUs attached to a machine, whereas Spark assumes homogenous cluster.
- Spark processes batches of data in the form of RDDs (useful for batch processing algos). TensorFlow can, in addition, efficiently handle one example flowing through the graph in one iteration.


Variation in test error for Synchronous SGD

The following result is plotted for a test set size of 15,000. The test error decreases to 28% in the first 100 iterations and then the decrease is very slow with some minor oscillations. It is around 25% at the 3000th iteration.




Variation in test error for Asynchronous SGD

The convergence graph is like synchronous SGD, with minor variations shown in the next section.




Difference between Asynchronous and Synchronous SGD

Asynchronous shows slightly faster convergence but has slightly more oscillations than synchronous in the later stages.



Batch SGD

The error convergence is very similar but Batch sizes of 100 and 1000 show a much smoother convergence graph than SGD.



Bottleneck for synchronous and asynchronous SGD

The bottleneck for both was found to be compute.

There are three major steps in SGD:

- 1) Read the data from disk
- 2) Compute local gradient
 - Gather the necessary indices from 'w' and compute dot product and local gradient
- 3) Update 'w'

It takes ~10-100ms for steps up to 2 and then ~1-2s for step 3. So, adding a sparse gradient vector to a dense vector 'w' was the bottleneck. [Note: `tf.gather(w, indices)` was placed at node 0 to reduce network load]

'Top' shows almost 100% CPU utilization at node 0, which is responsible for updating w.

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
29002	ubuntu	20	0	4025736	3.748g	17636	R	99.7	19.1	4:44.81	python
23000	ubuntu	20	0	11.953g	481120	1208	S	5.6	2.3	7:14.50	java
22846	ubuntu	20	0	10.954g	569568	1340	S	0.7	2.8	0:41.17	java
24193	ubuntu	20	0	1850336	359372	1780	S	0.7	1.7	1:12.76	java
9653	ubuntu	20	0	2607884	217108	0	S	0.3	1.1	3:02.02	java
10186	ubuntu	20	0	2607884	211364	0	S	0.3	1.0	3:03.29	java
12706	ubuntu	20	0	10.899g	135256	0	S	0.3	0.7	44:02.87	java

Disk and network utilizations were found to be very low. (Disk utilization could peak when checkpointing kicks in but this would be rare and not observed with initial iterations)

For 20 iterations, the number of bytes read and written were 163840 bytes/1642496 bytes and the NW read/write were 787630 bytes/778095 bytes.

The following figure shows it for one instance of time.

```
ubuntu @ ~
[9] → dstat -c --top-cpu -dn
----total-cpu-usage---- -most-expensive- -dsk/total- -net/total-
usr sys idl wai hiq siq|_cpu process_|_read_ writ_|_recv_ send_
 3  1  96  0  0  0|ksoftirqd/0 0.0| 32k 72k| 0 0
23 11 63  0  0  2|startserver.p 28|  0  0|1903B 1230B
25 11 63  0  0  1|startserver.p 29|  0  0|2576B 1296B
25 11 63  0  0  1|startserver.p 28|  0  0|6859B 5455B
21 11 68  0  0  1|startserver.p 26|  0 60k|1837B 987B
20 10 69  0  0  1|startserver.p 17|  0  0|2708B 1346B
19 13 68  0  0  0|asyncsgd.py 15|  0  0|6719B 6277B
20  9 70  0  0  1|startserver.p 26|  0  0|2557B 3210B
21  8 71  0  0  1|startserver.p 14|  0  0|5142B 2300B
21  9 69  0  0  1|startserver.p 16|  0 48k|7075B 5806B
```

Assignment 3 Part B

Group -11

Abhanshu Gupta (agupta89@wisc.edu)

Parikshit Sharma (psharma43@wisc.edu)

Vishnu Lokhande (lokhande@wisc.edu)

Application 1

Pagerank application was developed using graphX api's. Two key Api's from graphX which were used were (a) outerJoinVertices and (b) aggregateMessages.

This application has been compared with Assignment1-PartC-Question3 where just spark api's were used to run the pagerank algorithm.

Inorder to make a uniform comparison, the configurations parameters were set the same in both the applications. Specifically, these were the configuration parameters.

```
("spark.locality.wait", "0")  
("spark.driver.memory", "1g")  
("spark.executor.memory", "18g")  
("spark.executor.instances", "5")  
("spark.executor.cores", "4")  
("spark.task.cpus", "1")  
("spark.serializer", "org.apache.spark.serializer.KryoSerializer")  
("spark.logLineage", "false")
```

The output of the application is saved on hdfs at /pagerank_results. The application completion time is measured by the time taken by the application from the beginning of reading time to saving the results on hdfs. Here is summary.

	<i>Pagerank (spark api's)</i>	<i>Pagerank (graphX api's)</i>
<i>Time Taken</i>	29 min	5.4 min
<i>Network Read</i>	13242 MB	18344 MB
<i>Network Write</i>	13549 MB	19026 MB
<i>Storage Read</i>	1434 MB / 7 MBps	1368 MB / 9 MBps
<i>Storage Write</i>	24249 MB / 2 MBps	4807 MB / 2 MBps
<i>Number of Stages</i>	24	68
<i>Number of Tasks</i>	480	544

Comparison

Additional Benefits of graphX:

GraphX provides api's to work directly with the edge and vertex attributes of a graph. This makes developing a graph based application lot easier. One example of this is the "aggregateMessages" api, which passes messages from previous vertices to next vertices over the edges.

GraphX provided "TripletFields.Src" functionality, which helps populate only the necessary fields of an RDD/graph thereby further improving the efficiency.

GraphX when used for graph applications provides better visualizations in the form of triplet views, vertex views and edge views. This is not the case for spark as there is only RDD based view.

graphX vs. Spark

It is observed that pagerank using graphX takes lot less time compared to that of Spark. This is possibly because the ranks graph is being cached in every iteration in the graphX application unlike the spark application. Ranks are used in every iteration, hence, caching them seems to improve the performance a lot. We did not cache ranks in the spark application because we did not get much improvements then as we were dealing with small size datasets. Another possible reason might be that the outerJoinVertices api used in graphX might be more optimized than that of graphX. The number of tasks in spark are less than that of graphX. The number of tasks of an application usually dependent on the level of parallelism which is related to the number of partitions in the application. We used 20 partitions in spark app (because we used 20 partitions before) as compared to 8 partitions in graphX app (default for the datasize).

Application 2

Input: For this application input is the output file "Question2_words.txt" of PartB from Assignment2. We have written a python script "language_cleanup.py" to cleanup web addresses, hashtags and non-alphanumeric words and stored it in the new file "graph_input.txt". This file will serve as input for all sub-applications that are developed under Application2.

Question1

Implementation: PartBApplication2Question1.scala

Logic:

1. Read the "graph_input.txt" file and create a VertexRDD which is a RDD(VertexId,Array[String]) using 'map' function.
2. Create EdgeRDD by doing cartesian product of vertexRDD with itself and then filter out the entries where there is no common word between two vertices
3. Create a graph using VertexRDD and EdgeRDD.
4. Get the triplet view of the graph and use filter on the condition that source word list should be larger than the destination word list.
5. Get the count of the filtered RDD as answer

Output: For my case of 262 vertex the total number of edges that satisfy the criteria: 33648

Question2

Implementation: PartBApplication2Question2.scala

Logic:

1. Read the "graph_input.txt" file and create a VertexRDD which is a RDD(VertexId,Array[String]) using 'map' function.
2. Create EdgeRDD by doing cartesian product of vertexRDD with itself and then filter out the entries where there is no common word between two vertices.
3. Create a graph using VertexRDD and EdgeRDD.

4. Get the aggregated message of counter from each neighbor, thereby creating a neighbor RDD which stores the number of neighbors of each vertex.
5. Find the maximum degree in the created RDD.
6. Filter the RDD to only have entries corresponding to maximum degree.
7. If RDD has only one entry then the entry is the answer.
8. Else find the entry which has maximum words and that entry will be the answer.

Output: For my case of 262 vertex the vertex most popular corresponds to time interval: 53

Question3

Implementation: PartBApplication2Question3.scala

Logic:

1. Read the "graph_input.txt" file and create a VertexRDD which is a RDD(VertexId,Array[String]) using 'map' function.
2. Create EdgeRDD by doing cartesian product of vertexRDD with itself and then filter out the entries where there is no common word between two vertices.
3. Create a graph using VertexRDD and EdgeRDD.
4. Get the aggregated message of word size from each neighbor and reducing it by adding, thereby creating a neighbor RDD which stores the total number of words in neighbors of each vertex.
5. Use map function to take average value on each entry of RDD.
6. Print the RDD values.

Output: The output is list of tuples of the form (VertexId,Average Words in neighbor), printed one vertex per line.

Question4

Implementation: PartBApplication2Question4.scala

Logic:

1. Read the "graph_input.txt" file and create a VertexRDD which is a RDD(VertexId,Array[String]) using 'map' function.
2. Flatten the vertexRDD to get list of words and then reduce them on key to get words and their count.
3. Find the maximum count in the created RDD.
4. Print any entry randomly from the RDD that has its count equal to maximum count.

Output: For my case of 262 vertex the word most popular comes out to be: rt

Question5

Implementation: PartBApplication2Question5.scala

Logic:

1. Read the "graph_input.txt" file and create a VertexRDD which is a RDD(VertexId,Array[String]) using 'map' function.
2. Create EdgeRDD by doing cartesian product of vertexRDD with itself and then filter out the entries where there is no common word between two vertices.
3. Create a graph using VertexRDD and EdgeRDD.
4. Call the connectedComponents() algorithm on this graph to get connected components.

5. Reduce the RDD based on cluster's min vertex and use a counter to keep track of nodes.
6. Print the RDD value that has the maximum count as the biggest cluster.

Output: For my case of 262 vertex the biggest size of subgraph comes to be: 262

Question6

Implementation: PartBApplication2Question6.scala

Logic:

1. Read the "graph_input.txt" file and create a VertexRDD which is a RDD(VertexId,Array[String]) using 'map' function.
2. Flatten the vertexRDD to get list of words and then reduce them on key to get words and their count.
3. Find the maximum count in the created RDD.
4. Locate a word in the RDD that has maximum count.
5. Filter and print all the vertex in vertexRDD that have the popular word in its word list.

Output: For my case of 262 vertex the all the 262 vertex has the most popular word rt